



Name: Syeda Malika Zainab Kazmi

Batch: Pana-101-P011

Course: AI-101

REPORT WRITING

Task 1 — Lab Document Practice & Prompting Report

- **What I Learned from the Lab**

The lab is based on the "AI Prompting in 2026" crash course from Panaversity's Agent Factory. Going through it, the biggest realization was that AI models work like a stateless colleague they have zero memory between sessions and can only work with what's inside the current context window. This completely changed how I approach writing prompts. I used to type short one-liners like a Google search. After this lab, I understand that the quality of the output is almost entirely determined by how well you brief the AI upfront.

The 13 concepts covered were organized into four parts: how AI knows things, how to talk to AI well, working with multimodal inputs, and choosing the right tools safely. Each concept built on the previous one in a logical way.

- **How My Prompting Improved**

At the start, I was writing prompts like "**give me workout ideas**" or "**write a report on X.**" The outputs were generic and forgettable exactly what the lab calls "slop." After applying the lab concepts, my prompts started including role context, specific constraints, output format, and explicit feedback loops. For example, instead of "**write a blog post,**" I first asked for 3 outline options, picked the best one, asked for a bullet expansion, graded it out of 10, and only then asked for the full draft. The difference in output quality was significant.

- **Which Techniques Helped the Most**

The single most impactful technique was the **brainstorm-iterate loop** (Concept 7). Forcing the AI to give 3–5 options before expanding any of them, then giving explicit feedback ("reject option 2, keep option 4 but change X") pushed the AI past its default "average internet answer" and into something actually useful.

Close second was the **rubric-based critique** (Concept 6). Asking the AI to score my work 1–10 on specific criteria like clarity, structure, evidence and then asking what change would raise each score the most, completely eliminated the sycophantic "great job!" responses. Forcing a number forces honesty.

Third was **context front-loading** (Concept 4). I started attaching files, pasting relevant background, and stating constraints before asking my question. Same model, dramatically better answers. The analogy that stuck with me was: brief the AI like you'd brief a smart new colleague on their first day they're capable but they know nothing about your situation unless you tell them.

- **Challenges Faced and How I Solved Them**

The biggest challenge was resisting the urge to accept the first output. Habitually, I'd ask a question, get a fluent-sounding paragraph, and move on. The lab forced me to realize that fluent $\neq$ correct or useful. My fix was to always ask for a critique after getting an output, and to set a scoring target (like 9/10) before stopping.

Another challenge was context rot keeping one long conversation going across multiple unrelated topics, which caused the AI to start mixing up earlier constraints with new ones. The fix was simple: start a new conversation when the topic changes, and use Projects to carry persistent context without the noise.